

Q-Learning – How to?

Chair of Automation and Information Systems

Technical University of Munich

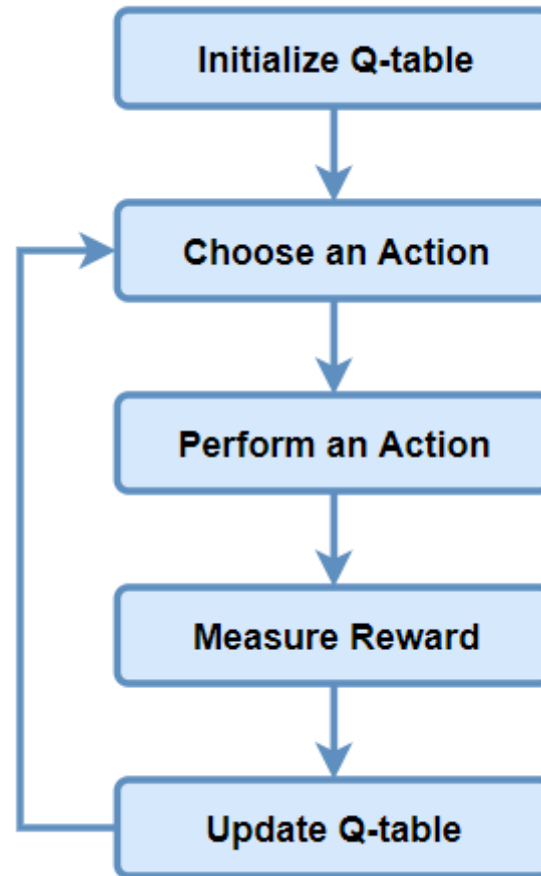


Topics

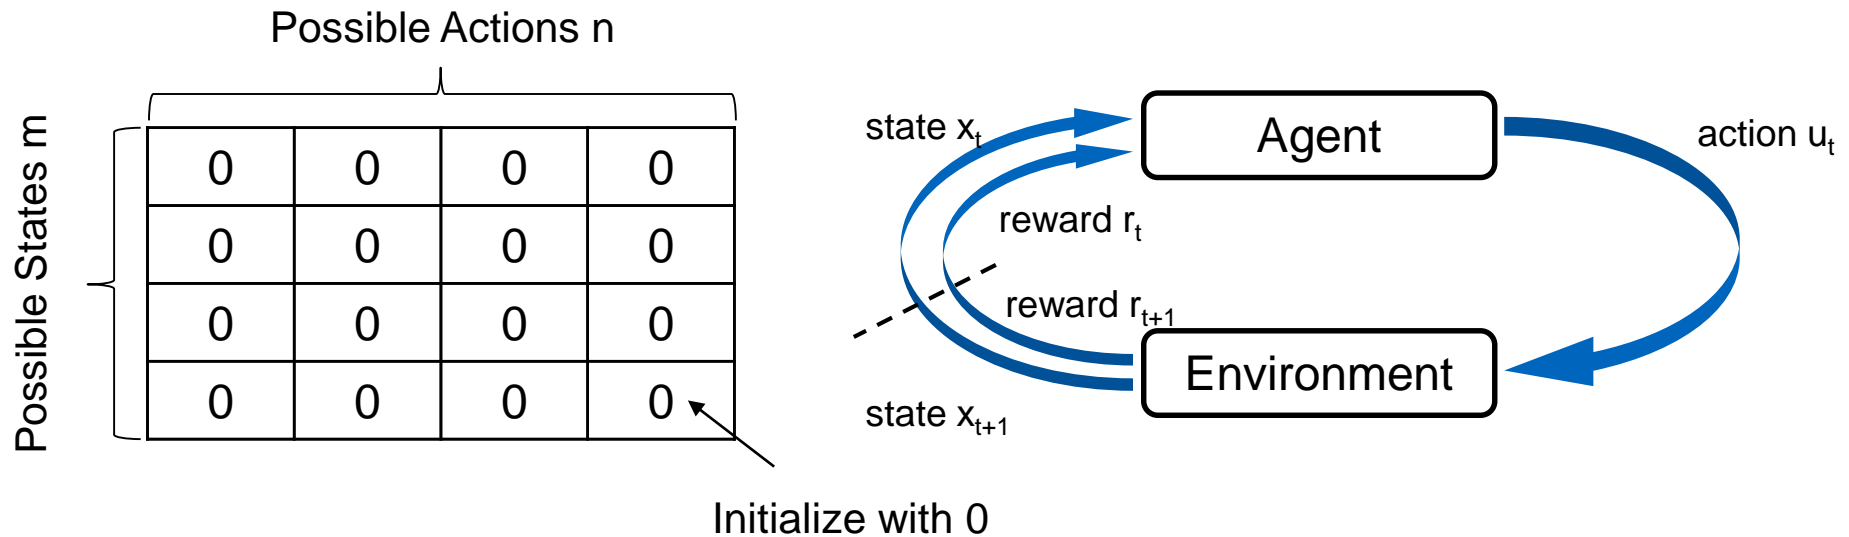
1. Q-Learning Process Steps
2. Exercise: Grid World
3. Exercise: Crane-Simulation

1. **Q-Learning Process Steps**
2. Exercise: Grid World
3. Exercise: Crane-Simulation

Q-Learning Process Steps



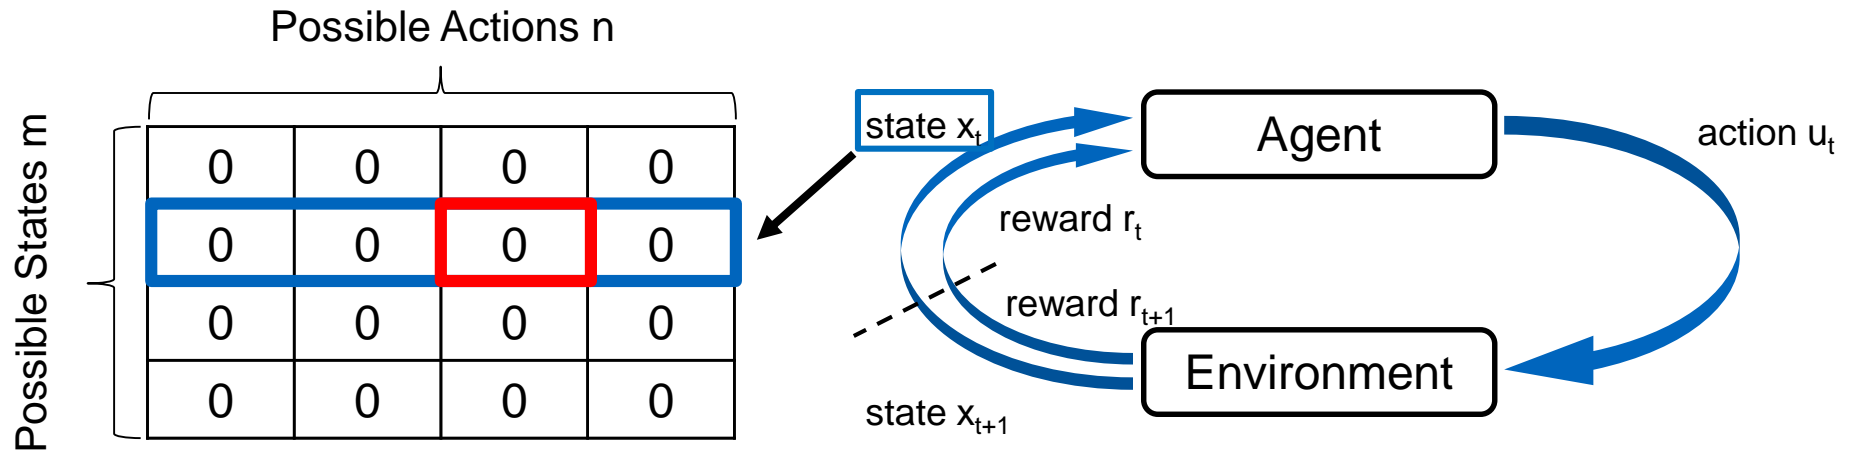
Q-Learning Process Steps



1. Initialize Q-Table

- Use discretization for environments with continuous state variables
- If the state is described by multiple state variables X_1, X_2, \dots, X_k (such as position, velocity and acceleration) with corresponding number of discretization intervals m_1, m_2, \dots, m_k , the Q-Table is of dimension $m_1 \times m_2 \times \dots \times m_k \times n$

Q-Learning Process Steps



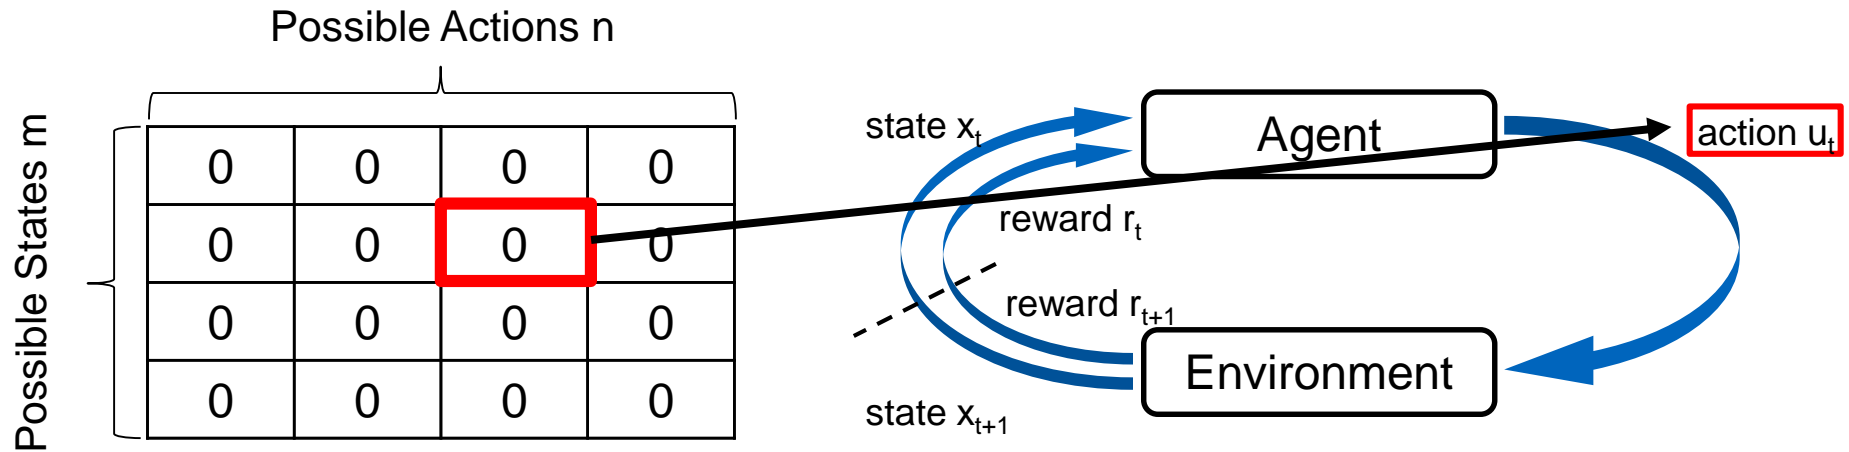
2. Choose an Action:

- Get the current state x_t
- Perform ϵ -greedy strategy:

$$\text{Action} = \begin{cases} \max_u Q(x, u), & R > \epsilon \\ \text{Random } u, & R \leq \epsilon \end{cases}, \text{ } R \text{ is uniform random value in } [0,1] \text{ and } \epsilon \in [0,1]$$

→ when highest Q-Value occurs more than once, choose randomly from actions with this value

Q-Learning Process Steps




3. Perform an Action

Q-Learning Process Steps

Possible Actions n

Possible States m

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0



4. Measure reward


And remember new state x_{t+1}

Q-Learning Process Steps

Possible Actions n

Possible States m

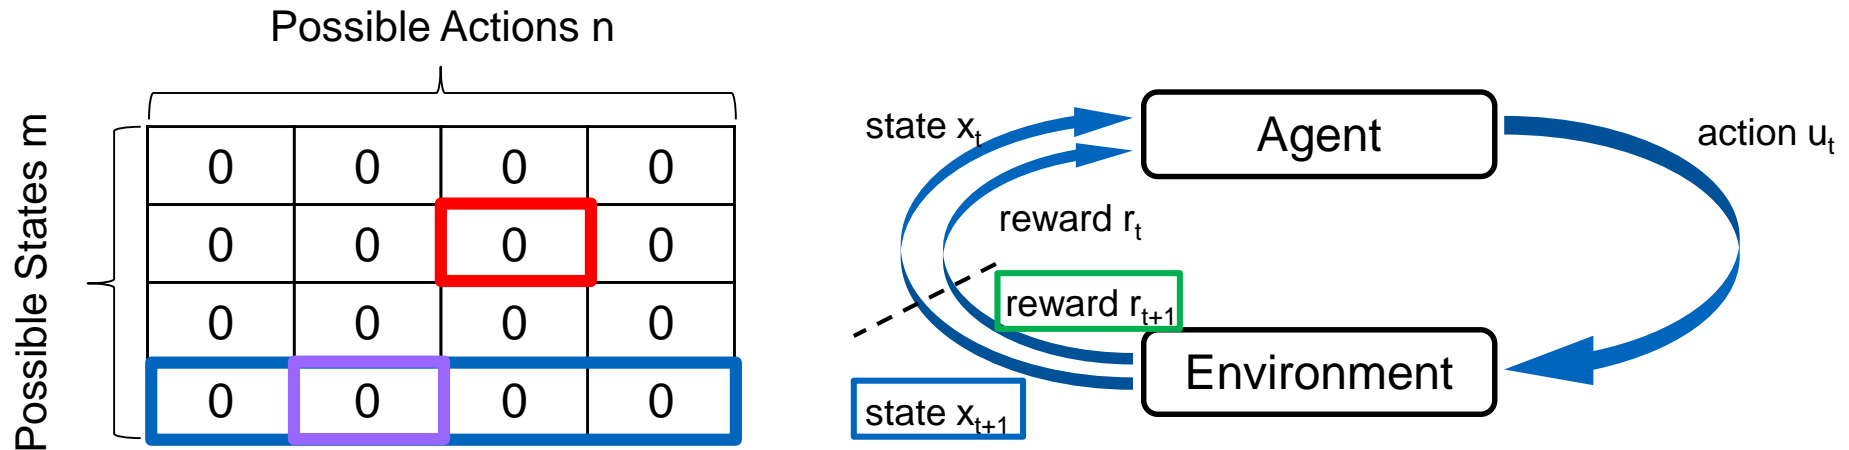
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0



5. Update Q-Table

- Go to new state x_{t+1}
- Find the highest possible Q-value of state x_{t+1} (Q-value = expected future reward)

Q-Learning Process Steps



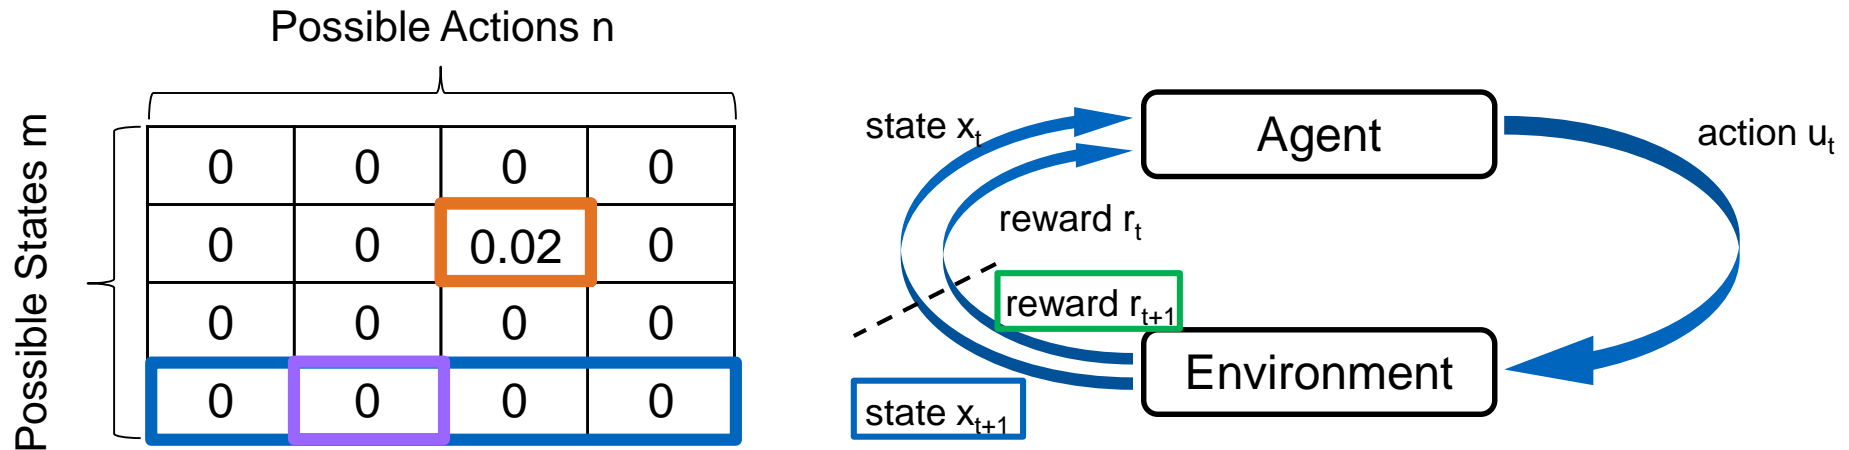
5. Update Q-Table

- Go to new state x_{t+1}
- Find the highest possible Q-value of state x_{t+1} (Q-value = expected future reward)
- Update Q-value for the original state-action pair:

$$Q_{k+1}^{\pi}(x_t, u_t) = (1 - \alpha) \cdot Q_k^{\pi}(x_t, u_t) + \alpha \cdot (r_{t+1} + \gamma \cdot \arg\max_u Q(x_{t+1}, u))$$

$r_{t+1} = 20, \quad \alpha = 10^{-3}, \quad \gamma = 0.95$

Q-Learning Process Steps



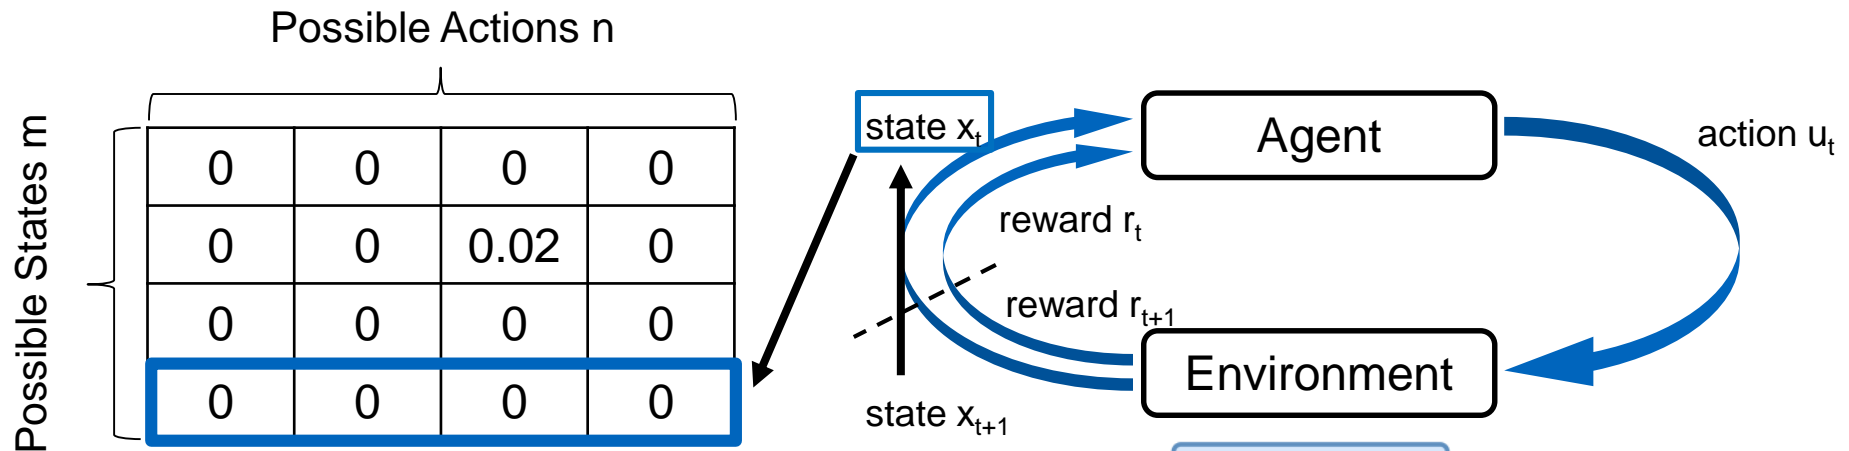
5. Update Q-Table

- Go to new state x_{t+1}
- Find the highest possible Q-value of state x_{t+1} (Q-value = expected future reward)
- Update Q-value for the original state-action pair:

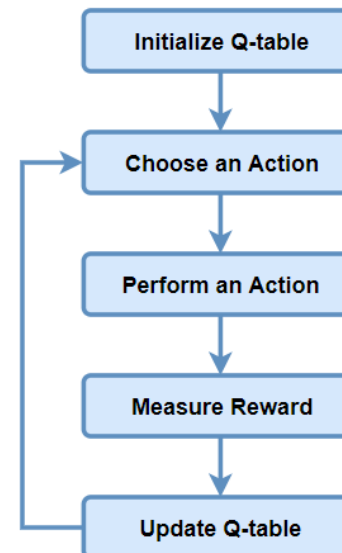
$$Q_{k+1}^{\pi}(x_t, u_t) = (1 - \alpha) \cdot Q_k^{\pi}(x_t, u_t) + \alpha \cdot (r_{t+1} + \gamma \cdot \arg\max_u Q(x_{t+1}, u)) = 0.02$$

$r_{t+1} = 20, \quad \alpha = 10^{-3}, \quad \gamma = 0.95$

Q-Learning Process Steps




Go back to step 2 and continue



1. Q-Learning Process Steps
2. **Exercise: Grid World**
3. Exercise: Crane-Simulation

Grid World – Task 1




- Small Grid World: 2x2
- Go from a starting point (S) to a target location (T)
- Possible actions: up, down, left or right (crashing into the boundary → position remains the same, but counts as new step for reward calculation)
- Rewards:

x_{t+1}	$r_{t+1}(x_{t+1})$
E	-1
T	+100
B	-0.5
S	-1

- Task 1:
 - Perform the first iteration of Q-Learning (Steps 1-5) by hand
 - Use learning rate $\alpha = 0.1$ and discount factor $\gamma = 1$
 - Neglect epsilon-greedy strategy

Grid World – Task 2




– Task 2:

- After 5 completed episodes with epsilon-greedy strategy ($\epsilon = 0.05$) the Q-Table looks like shown below
- Compute the first iteration of Q-Learning for episode 6, starting again at (S) by hand
- As before: Neglect epsilon-greedy strategy

	Up	Down	Left	Right
E	0	-0.1	-0.1	0
T	0	0	0	0
S	-0.1	-0.1	-0.1	7.91
B	40.95	0	-0.1	0


Grid World – Task 3



- Imagine an epsilon-greedy strategy where ϵ is computed as a function of the current episode:
→ $\epsilon(\text{episode}) = 0.9999^{\text{episode}}$
- Epsilon gets very small over time
- Task 3: What does the Q-Table look like after an infinite amount of episodes?

$$Q_{k+1}^{\pi}(x_t, u_t) = (1 - \alpha) \cdot Q_k^{\pi}(x_t, u_t) + \alpha \cdot \left(r_{t+1} + \gamma \cdot \underset{u}{\operatorname{argmax}} Q(x_{t+1}, u) \right)$$

Grid World – Task 4



- Task 4: What happens, if we apply the following reward structure to the Grid World problem?

x_{t+1}	$r(x_{t+1})$
E	-1
T	+100
B	+1
S	-1

1. Q-Learning Process Steps
2. Exercise: Grid World
3. **Exercise: Crane-Simulation**

Crane Simulation


Source: Chapter Machine Learning, Lecture ISMLP, AIS 2021

- Quadratic Warehouse \rightarrow discretized into a 10 x 10 grid
- A product has to be transported with a crane from a starting point (S) to a target location (T)
- When the crane crashes into an high rack storage (O), the game ends
- Possible actions: up, down, left or right (crashing into the boundary \rightarrow position remains the same)
- Rewards:

x_{t+1}	$r_{t+1}(x_{t+1})$
White field / (S)	-1
(T)	+100
(O)	-100



Top view:



Crane Simulation – Tasks

Tasks:

- Open *crane_simulation_Template.py*
- In class *CraneSim*: define the grid and all required variables
- In class *Q_Agent*: implement the epsilon-greedy policy and the update function for the Q-Table
- In function *play*: implement all necessary steps to run one episode of the simulation

